

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, accuracy_score
```

```
df = pd.read_csv('/content/archive (3).zip')
```

```
df.head()
```

	Age	Sex	Chest pain type	BP	Cholesterol	FBS over 120	EKG results	Max HR	Exercise angina	ST depression	Slope of ST	Number of vessels fluro	Thallium	Heart Disease
0	70	1	4	130	322	0	2	109	0	2.4	2	3	3	Presence
1	67	0	3	115	564	0	2	160	0	1.6	2	0	7	Absence
2	57	1	2	124	261	0	0	141	0	0.3	1	0	7	Presence
3	64	1	4	128	263	0	0	105	1	0.2	2	1	7	Absence

Next steps:

Generate code with df

View recommended plots

New interactive sheet


```
df.tail()
```

	Age	Sex	Chest pain type	BP	Cholesterol	FBS over 120	EKG results	Max HR	Exercise angina	ST depression	Slope of ST	Number of vessels fluro	Thallium	Heart Disease
265	52	1	3	172	199	1	0	162	0	0.5	1	0	7	Absence
266	44	1	2	120	263	0	0	173	0	0.0	1	0	7	Absence
267	56	0	2	140	294	0	2	153	0	1.3	2	0	3	Absence
268	57	1	4	140	192	0	0	148	0	0.4	2	0	6	Absence

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 270 entries, 0 to 269
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                    270 non-null   int64
1   Sex                    270 non-null   int64
2   Chest pain type        270 non-null   int64
3   BP                     270 non-null   int64
4   Cholesterol            270 non-null   int64
5   FBS over 120           270 non-null   int64
6   EKG results            270 non-null   int64
7   Max HR                 270 non-null   int64
8   Exercise angina        270 non-null   int64
9   ST depression          270 non-null   float64
10  Slope of ST            270 non-null   int64
11  Number of vessels fluro 270 non-null   int64
12  Thallium                270 non-null   int64
13  Heart Disease          270 non-null   object
dtypes: float64(1), int64(12), object(1)
memory usage: 29.7+ KB
```

```
df.isnull().sum()
```



	0
Age	0
Sex	0
Chest pain type	0
BP	0
Cholesterol	0
FBS over 120	0
EKG results	0
Max HR	0
Exercise angina	0
ST depression	0
Slope of ST	0
Number of vessels fluoro	0
Thallium	0
Heart Disease	0

df.head(4)


```
label_encoder = LabelEncoder()
df['Heart Disease'] = label_encoder.fit_transform(df['Heart Disease'])
```

```
X = df.drop('Heart Disease', axis=1)
y = df['Heart Disease']
```



```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

```
model = RandomForestClassifier(random_state=42)
```

```
model.fit(X_train, y_train)
```



RandomForestClassifier





```
RandomForestClassifier(random_state=42)
```

```
y_pred = model.predict(X_test)
```

```
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred, target_names=label_encoder.classes_)
```


```
print(f"Accuracy: {accuracy:.2f}")
```



```
Accuracy: 0.80
```

```
print("Classification Report:")
```

```
print(report)
```



```
Classification Report:
```

	precision	recall	f1-score	support
Absence	0.81	0.88	0.84	33
Presence	0.78	0.67	0.72	21
accuracy			0.80	54
macro avg	0.79	0.77	0.78	54
weighted avg	0.79	0.80	0.79	54

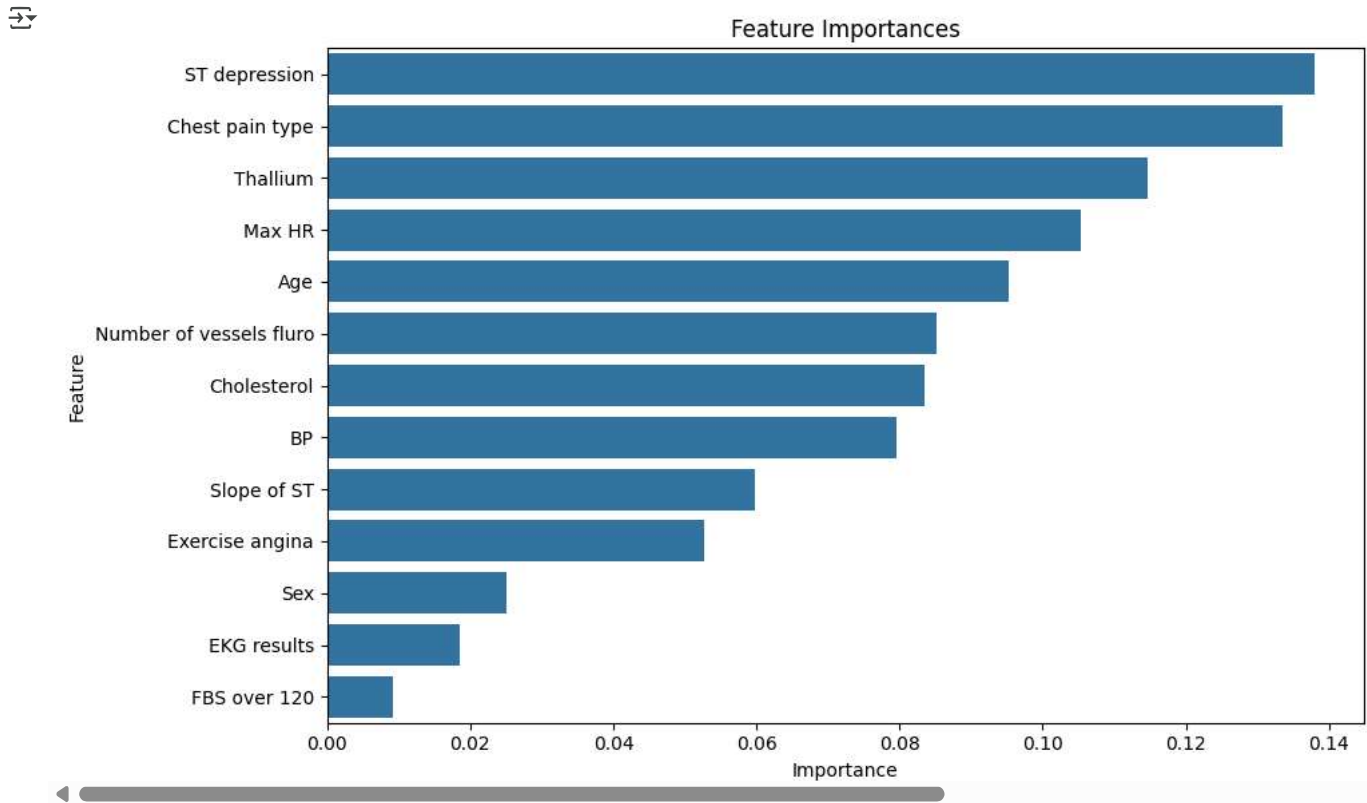
```
importances = model.feature_importances_
```

```

features = X.columns
indices = importances.argsort()[::-1]

plt.figure(figsize=(10, 6))
sns.barplot(x=importances[indices], y=features[indices])
plt.title("Feature Importances")
plt.xlabel("Importance")
plt.ylabel("Feature")
plt.tight_layout()
plt.show()

```



```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from xgboost import XGBClassifier
from sklearn.metrics import classification_report, accuracy_score

df = pd.read_csv('/content/archive (3).zip')

df['Heart Disease'] = df['Heart Disease'].map({'Presence': 1, 'Absence': 0})

label_encoder = LabelEncoder()
df['Heart Disease'] = label_encoder.fit_transform(df['Heart Disease'])

X = df.drop('Heart Disease', axis=1)
y = df['Heart Disease']

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X = df.drop('Heart Disease', axis=1)
y = df['Heart Disease']

X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42
)

```

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

```
# Drop rows where the target is NaN
df = df.dropna(subset=['Heart Disease'])
```

```
# Proceed with feature/target split
X = df.drop('Heart Disease', axis=1)
y = df['Heart Disease']
```

```
# (Continue with StandardScaler, train_test_split, etc.)
```

```
print(df['Heart Disease'].isna().sum())
```

```
0
```

```
print(df['Heart Disease'].isna().sum())
```

```
0
```

```
df = df.fillna(df.mean(numeric_only=True))
```

```
log_reg = LogisticRegression()
```

```
log_reg.fit(X_train, y_train)
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
    LogisticRegression
```

```
LogisticRegression()
```

```
log_preds = log_reg.predict(X_test)
```

```
print("=== Logistic Regression Results ===")
```

```
print("Accuracy:", accuracy_score(y_test, log_preds))
```

```
print(classification_report(y_test, log_preds))
```

```
=== Logistic Regression Results ===
```

```
Accuracy: 0.9074074074074074
```

	precision	recall	f1-score	support
0	0.89	0.97	0.93	33
1	0.94	0.81	0.87	21
accuracy			0.91	54
macro avg	0.92	0.89	0.90	54
weighted avg	0.91	0.91	0.91	54

```
xgb = XGBClassifier(use_label_encoder=False, eval_metric='logloss')
xgb.fit(X_train, y_train)
```

```

/usr/local/lib/python3.11/dist-packages/xgboost/core.py:158: UserWarning: [14:29:40] WARNING: /workspace/src/learner.cc:740:
Parameters: { "use_label_encoder" } are not used.

```

```
warnings.warn(msg, UserWarning)
```

```

XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric='logloss',
               feature_types=None, gamma=None, grow_policy=None,
               importance_type=None, interaction_constraints=None,
               learning_rate=None, max_bin=None, max_cat_threshold=None,
               max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
               max_leaves=None, min_child_weight=None, missing=nan,
               monotone_constraints=None, multi_strategy=None, n_estimators=None,
               n_jobs=None, num_parallel_tree=None, random_state=None, ...)

```

```
xgb_preds = xgb.predict(X_test)
```

```

print("=== XGBoost Results ===")
print("Accuracy:", accuracy_score(y_test, xgb_preds))
print(classification_report(y_test, xgb_preds))

```

```

=== XGBoost Results ===
Accuracy: 0.8148148148148148

```

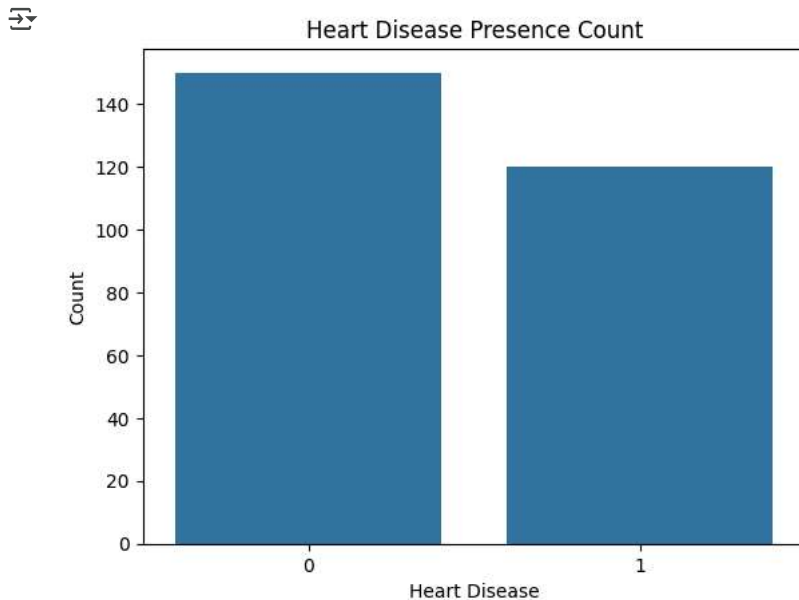
	precision	recall	f1-score	support
0	0.83	0.88	0.85	33
1	0.79	0.71	0.75	21
accuracy			0.81	54
macro avg	0.81	0.80	0.80	54
weighted avg	0.81	0.81	0.81	54

```

import seaborn as sns
import matplotlib.pyplot as plt

sns.countplot(data=df, x='Heart Disease')
plt.title('Heart Disease Presence Count')
plt.xlabel('Heart Disease')
plt.ylabel('Count')
plt.show()


```




```

sns.histplot(data=df, x='Age', hue='Heart Disease', kde=True, bins=20)
plt.title('Age Distribution by Heart Disease')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()

```



```
sns.boxplot(data=df, x='Chest pain type', y='Age', hue='Heart Disease')
plt.title('Chest Pain Type vs Age by Heart Disease')
plt.xlabel('Chest Pain Type')
plt.ylabel('Age')
plt.show()
```



```
plt.figure(figsize=(12, 8))
corr = df.corr(numeric_only=True)
sns.heatmap(corr, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Feature Correlation Heatmap')
plt.show()
```

